

Analysis of stock market recommendations using computer vision

Submitted in fulfilment of the requirements
of the degree of

Masters of Technology

by

Pronoy Mandal (Registration no. 20 - 27 - 09)

Supervisor:

Dr. (Mr.) Dasari Srikanth

Co - supervisor:

Mr. Himanshu Chaudhary



Department of Applied Mathematics
Defence Institute of Advanced Technology (Pune)
2021 - 22

CERTIFICATE

This is to certify that the project entitled “**Analysis of stock market recommendations using computer vision**” is a genuine work of **Pronoy Mandal (Registration no. 20 - 27 - 09)** submitted to the Defence Institute of Advanced Technology (Pune) in fulfilment of the requirement for the award of **Masters of Technology in Data Science**.

Certified: _____

Dr. (Mr.) Dasari Srikanth
Supervisor/guide

Certified: _____

Mr. Himanshu Chaudhary
Co - supervisor
Assistant General Manager, Integrated Surveillance-
Department, SEBI

Certified: _____

Dr. (Mr.) Somanchi VSSNVG Krishnamurthy
Head of the Department of Applied Mathematics

Dissertation Approval for M.Tech.

This project report entitled **Analysis of stock market recommendations using computer vision** by **Pronoy Mandal** is certified for the degree of **Masters of Technology in Data Science**.

Examiners

1. _____

2. _____

Date:

Place:

DECLARATION

I declare that this written submission represents my ideas in my own words, and where other ideas or words have been included, I have adequately cited and referenced the original sources. I also declare that I have adhered to all principles of academic honesty and integrity and that I have not misrepresented, fabricated, or falsified any idea, data, fact or source in my submission. I understand that any violation of the above will be cause for disciplinary action by the institute and can also evoke penal action from the sources which have thus not been properly cited or from whom proper permission has not been taken when needed.

Pronoy Mandal (Registration no. 20 - 27 - 09)

Date:

Abstract

There are a large number of business NEWS channels broadcasting a variety of topics throughout morning and evening primetime. These range from simple stock market recommendations (strategies to BUY/SELL/HOLD) to detailed fundamental and technical analyses of various companies to in-depth quarterly performance reviews of various parties in the market. With the availability of such a volume of information from public sources; it becomes difficult to assimilate all the information in a single place and analyse it for meaningful information.

This project encompasses the ability to analyse NEWS which gives stock market recommendations (a subset of business NEWS) and presents the results in a tabular format to the user for analysis. The tabular format includes basic information from the broadcast such as the telecast date of the show, the analyst presenting the same, the stock which is being discussed, the NSE listing symbol of the concerned stock, the recommendation etc. (totalling 8 parameters). The exact reason why the Securities and Exchange Board of India (SEBI) is attempting to analyse NEWS videos is something that can't be specified within the scope of this report due to its confidentiality. However, it can be said that this project's output forms the base as well as the input of many anomaly (in the context of violations in the securities market) detection models at a later stage.

This project goes a step ahead to ease and automate the job of SEBI officers from the surveillance department by deploying the same onto an MLOPs platform and creating an equivalent free-running workflow in a virtual environment.

The aforementioned reason forms the principal motivation for doing this project i.e. it helps automate a large part of a task that is prone to human errors and mistakes such as collecting data manually by watching the NEWS shows and at the same time reduces the time required to carry out such tasks. Apart from that, since an MLOPs workflow is included in the project, this project helps establish various principles which need to be followed for any ML project in general so that it can be deployed into production.

Table of Contents

1	Introduction	1
1.1	Dissertation name and title	1
1.2	Problem statement	1
1.3	Proposed solution	1
1.4	Scope and report contents	2
2	Literature review	4
2.1	YoLo – A brief discussion on v3 and v4 for real-time object detection . . .	4
2.1.1	YoLov3	5
2.1.2	YoLov4 – improvements over YoLov3	7
2.2	Performance metrics for evaluating object detection models	7
2.3	Guidelines for a good YoLo project	7
2.3.1	For training	7
2.3.2	For testing or detection	8
3	Methodology and development	9
4	Deployment in production	10
5	Results and discussions	11
6	Conclusions	12
	Bibliography	13

Acknowledgements**13**

List of Figures

2.1	A sample Darknet run on an image	5
2.2	An illustration showing all the important bounding box dimesnions being used in YoLov3	7

List of Tables

Chapter 1

Introduction

1.1 Dissertation name and title

My dissertation is titled **Analysis of stock market recommendations using computer vision**.

1.2 Problem statement

There is a large no. of daily broadcasts which take place throughout primetime (morning and evening) as well as outside of it. This project aims to digitize as well as summarise the recommendations cumulatively and present them to the end-user at the same time this project and its associated machine learning system are input to more sophisticated models for anomaly detection (in the context of violations in the securities market) at a later stage.

1.3 Proposed solution

The proposed solution is a deep learning model aimed at solving two fundamental machine learning problems at hand:

- (a) Object detection – Detect and classify important regions of interest in the frames of various NEWS broadcasts (or NEWS shows) i.e. regions having (but not restricted to) the name of the analyst/presenter, the recommendation being given, the company, commodity or the market segment being discussed and its corresponding listing symbol (if any), the various metrics of the recommendation such as stop loss, target price etc.
- (b) Optical Character Recognition (OCR) – Bilingual (English and Hindi) text obtained from the above regions of interest is passed onto sophisticated OCR models which have a text detection model and an actual language model(s) (corresponding to

the concerned languages) working under the hood to obtain or extract the required information.

The final output consists of a CSV (or) an excel file as required which summarises the above-obtained information.

1.4 Scope and report contents

Chapter 2 goes into a detailed discussion of the literature which has been studied while preparing this report. It includes several details such as (but not restricted to) how **YoLo** models or in general object detection models are evaluated and which performance metrics and common datasets are referred to for their speed and accuracy. It has additional details regarding how advanced deep learning models like **YoLov4** can be optimised for special use cases wherein necessary. The work presented in this dissertation may be completed sequentially or parallelly and is described as follows. Additionally, the chapters to which they pertain to are also mentioned.

- (a) Machine learning and computer vision: There are a total of 8 NEWS shows which are currently being targeted under this project. Machine learning used within this project is further divided into two parts: the one dealing with the recognition of various numerical parameters from the video frames of a broadcast (wherein active work is going on to increase efficiency as well as the accuracy of models) are random forest classifiers which have been trained on a large number of video frames which were earlier trained on **YoLov3** output (but needed additional training due to lesser accuracy) and templates followed by a much more robust and reliable framework i.e. YoLov4 which has been trained only for one NEWS show (i.e. *PehlaSaada* from CNBC Awaaz).

The second part of ML goes into OCR models wherein **no active work is being pursued** rather readymade models and OCR engines that are doing the job within tolerable error limits are being used. Both Tesseract – OCR and EasyOCR provided models have been used for this purpose and their performance metrics analysed. It should be noted that **no part of the report** goes into the details of the inner workings of the OCR models which are deep learning models in themselves.

As opposed to standard ML projects wherein a considerable amount of time is used for training a particular model; this project uses more time and space to carry out a variety of image processing operations on the captured video frames. This includes histogram equalization, simple grey level transformations, and complex erosion and dilation processes. After carrying out sufficient thresholding operations, these are fed into the ML models described above.

The details for the above i.e. the process of collecting the data, preparing the data for training and testing and using appropriate deep learning models for the same are discussed in chapter 3 in detail.

- (b) Deployment into production: Rarely do data science projects in academia ever reach production and into a streamlined architecture. There are several reasons which

prevent it from happening so. Such details, as well as reasons, have been given in brief in chapter 4's opening sections.

Fortunately, for this project, the deployment of the data models falls within its current scope. All the ML models which would be trained and tested would be loaded onto the proprietary **Ezmeral MLOPs** platform by **HPE**. The servers for the same reside alongside the required hardware resources on SEBI's internal servers. Doing this (without going into the details) would make the entire process streamlined i.e. the process of uploading videos, and training models on new data as soon as it is available and at the same time would decrease the manual work involved while doing so. As soon as the deployment has been completed, whether for all the NEWS shows or even a few of them, the internal server would be able to serve requests from within the organization domain (*POST* methods with *JSON* body) and at the same time would be able to abstract all the complex code and scripts running in the background.

Apart from the MLOPs deployment, it is more desirable for SEBI to have a continuously running workflow that shall consist of processing a bulk or a batch of videos at once. A workflow about this has been deployed on a standalone virtual machine as a part of this project. Which projects require an MLOPs workflow and which don't, as a result, have also been discussed in chapter 4.

Chapter 5 would then go on to discuss the various results which have been obtained at the end in terms of speed and accuracy of the YoLo models involved which metrics are suitable for evaluating projects spanning multiple ML models. Chapter 6 of the report would then finally present important conclusions drawn from the entire project.

Chapter 2

Literature review

2.1 YoLo – A brief discussion on v3 and v4 for real-time object detection

YoLo standing for You Only Look Once is(are) a set of advanced deep learning models for real-time object detection. The focal point for YoLo is that it manages to perform bounding box regression and classification at the same time. The architecture or methodology which separates YoLo and other object detection frameworks is that other sophisticated frameworks (RCNN and even Fast RCNN) attempt to repurpose existing classifiers and localisers to perform object detection. They apply a single model to multiple locations within an image with variations in scale and other properties of the image multiple times. High scoring regions are labelled and classified accordingly.

YoLo frameworks take a drastically different approach: they divide the region into a large number of regions (or boxes) and simultaneously apply a single convolutional neural network to the entire image. The bounding boxes and associated probabilities are received at the output. These regions are then weighted according to these probabilities. The reason why this method supersedes other popular methods is that other neural network-based deep learning architectures (such as R-CNN) may end up using multiple network evaluations (in the order of thousands) on a single image that takes up a considerable amount of processing time. Additionally, YoLo models look at the entirety of the image during the detection phase because of which their predictions are based on the global context of the image. Although the differences don't end here, these fundamental differences are enough to give YoLo a second to third order speed advantage over Fast R-CNN and R-CNN models.

The neural network for YoLo along with its entire compilation system is known as Darknet. A typical darknet run on an image, set of images, video or a real-time stream of image content (such as through a webcam) leads to the following three things as output:


- The detected classes of objects (e.g. horse, car etc.)
- The prediction confidence (e.g. 0.91, 0.86 etc.) and
- The time it took to carry out the prediction or detection (e.g. 6s, 12s etc.)

- Sophisticated high-level libraries will also produce the bounding boxes corresponding to each

Following is the output when a high-level library like [yolov4](#) is used to process an image as shown below.



(a) A sample image taken from the original Darknet [github repository](#)



(b) Bounding boxes and detected objects in the image

```
YOLOv4: Inference is finished
[array([[ 0.7864571,  0.55349463,  0.32891354,  0.4883267, 17.        ,
          0.90753984],
        [ 0.20731689,  0.72633255,  0.21659599,  0.22171444, 16.        ,
          0.9931677 ],
        [ 0.36157072,  0.5735911 ,  0.135803  ,  0.6407941 ,  0.        ,
          0.9047223 ]], dtype=float32)]
[Finished in 7.416s]
```

(c) The bounding box coordinates obtained

Figure 2.1: A sample Darknet run on an image

2.1.1 YoLov3

Joseph Redmond is credited with designing the original neural network architecture called Darknet. This was made using all low-level languages so that it could be made as flexible

as it can be. This architecture went on to produce a series of computer vision wonders in the field of object detection named as YoLo: the original one, v2, v3 and then v4. Even YoLov5 has been developed (the first one to be developed by an organisation as opposed to an individual) at the time of writing this report. However, we won't be discussing it here since it hasn't been used in the concerned project.

YoLov3 is the first version using an objectivity score for the prediction of bounding boxes. It also proceeds to add further connections to the backbone of the entire architecture. Additionally, detections are also carried out at three levels of granularity which greatly increased the inferencing accuracy on smaller objects as well as objects/ smaller objects placed close to each other. There are some important parts or sections in YoLov3 which proves its robustness to a variety of real-life object detection scenarios. They relate to better bounding box prediction, class prediction, their feature extractor as well better prediction across multiple scales provided in the input. We would be discussing each of them in brief in the following paragraphs.

Bounding box prediction

Any bounding box in a prediction is represented by a set of four parameters namely t_x and t_y : coordinates of the centre of the bounding box and its width and height represented by t_w and t_h respectively. In the event of the top left corner of the bounding box being displaced by (c_x, c_y) respectively, the following overall changes are calculated in the dimensions (assume p_w and p_h are the width and height of the bounding box prior).

$$\begin{aligned} b_x &= \sigma(t_x) + c_x \\ b_y &= \sigma(t_y) + c_y \\ b_w &= p_w e^{t_w} \\ b_h &= p_h e^{t_h} \end{aligned}$$

Unlike previous systems (or previous YoLo versions), the training phase assigns an objectness score to each bounding box. This score equals one if the current bounding box overlaps a ground truth object to an extent greater than any other previously obtained bounding box. If any previously obtained bounding box doesn't overlap the ground truth to a greater extent or overlaps only by a certain threshold (say 0.5 as stated in this paper) then the objectness score is simply zero. Following is a representation of all the important dimensions which are involved in this process.

Class prediction

SoftMax is one of the most common types of methods used for multi-label classification, however, the method has been leading to decrementing results for the class prediction accuracy in YoLov3 as well as has been unnecessary. Therefore, simple multiple logistic classifiers are being used for the purpose and binary cross-entropy has been used as a loss function. This different approach helped in the case of much larger and more sophisticated datasets such as the Open Images Dataset wherein labels may overlap frequently. SoftMax algorithm assumes that a single bounding box corresponds to exactly one label which is often not the case.




Figure 2.2: An illustration showing all the important bounding box dimensions being used in YOLOv3

Feature extraction

Feature extraction in YOLOv3 has been a hybridisation of the one in YOLOv2, Darknet – 19 (19 convolutional layers) and another residual network. But the most important part of this network is shortcut connections which increase the size of the network significantly, however, continue to supersede ResNet variants in terms of efficiency. Since this leads to a total of **53** convolutional layers it's named Darknet-53.

2.1.2 YOLOv4 – improvements over YOLOv3

2.2 Performance metrics for evaluating object detection models

2.3 Guidelines for a good YOLO project

The following rules for a good project using YOLO are mere *thumb – rules* and are not some strict guidelines to be followed and should be evaluated on a case to case basis for every project. Additionally, it should be noted that such rules may not be applicable for implementation in every object detection project. The following rules are bifurcated into those undertaken during the training and those in the testing (detection) phase.

2.3.1 For training

- (a) The values of `random` should be set to 1 in your `.cfg` file which allows training for multiple image or video resolutions simultaneously.

- (b) Every distinct object that is liable for detection must have an appropriate label in the dataset.
- (c) Precision may be increased by keeping the height and width of images or video frames as a multiple of 32.
- (d) The training dataset should be such that every object to be detected corresponds to an exactly similar object in the training dataset. Similarity should be in terms of size, no. of classes to be detected (c), overall spatial orientation, overall illumination, augmentation (if any) etc.

If the size of the training dataset is N and the no. of classes is stated as above then training must run for at least Nc iterations.

- (e) Training datasets should have as many **positive** examples as there are **negative** examples (i.e. images which don't have any object to be detected). Such negative examples shall return no bounding boxes when the detection or testing run is executed. These ensure an equal sensitivity of the model to both types of images as well as eliminate a lot of post-processing operations.
- (f) Sometimes it is desirable to run your detections with the `-show_imgs` option at the end so that it can be manually verified whether the predicted bounding boxes are correct or not. Seeing the detections and detecting some anomaly could be a direct implication of training runs going wrong or some inherent problem in the dataset.

2.3.2 For testing or detection

- (a) Increase network resolution in the same way as mentioned in 1. c.
- (b) It should be noted that retraining is not required in the event of loss in your dataset or any other unintended corruption. Only the *darknet* command should be available which can be used to perform detections using the pre-trained `.weights` file which was trained on the 416×416 resolution images.
- (c) To further enhance the accuracy, dataset training must proceed onto higher multiples of 32 such as 608×608 or 832×832 . In the event of a memory overflow (Out of memory), subdivisions in your `.cfg` file must be increased from 16 to 32 to 64 and so on.

Chapter 3

Methodology and development

Chapter 4

Deployment in production

Chapter 5

Results and discussions

Chapter 6

Conclusions

Acknowledgements

I am thankful to my institute Defence Institute of Advanced Technology (Pune) for considering my dissertation and providing help at all stages necessary during my work collecting information on my dissertation.

It gives me immense pleasure to express my deep and sincere gratitude to Dr.(Mr.) Dasari Srikanth (Dissertation guide) for his kind help and valuable advice during the development of the dissertation synopsis and his guidance and suggestions. Apart from that, I would like to thank Mr Himanshu Chaudhary (co-guide), Assistant General Manager from the Integrated Surveillance Department (ISD) of the Securities and Exchange Board of India (SEBI) for guiding me throughout this project and giving his valuable suggestions wherein necessary. I would like to express my deepest sense of gratitude to the employees of the Integrated Surveillance Department (ISD) and Information Technology Department (ITD) of SEBI who have put in their time and efforts to give additional guidance and suggestions during the project.

I would also like to thank the employees of Hewlett Packard Enterprise (HPE), a contractor of SEBI, who have allowed me to access their development environments, sophisticated hardware and cloud platforms, and associated tools and have been an immense help throughout this project.

I am deeply indebted to the Head of the Applied Mathematics Department, Dr. (Mr.) Somanchi VSSNVG Krishnamurthy, for giving me this valuable opportunity to do this project. I express my hearty thanks to him for his assistance, without which it would have been difficult to finish this project and review the project successfully. I would like to convey my deep sense of gratitude to all the teaching and non-teaching staff for their constant encouragement, support, and selfless help throughout the dissertation work. It is my great pleasure to acknowledge the help and suggestions that I received from the Department of Applied Mathematics.